

## Q&A Crib — If They Ask, You Say

Companion to the presenter teleprompt. Bold openers are the sentence to lead with; the rest is your backing. Grey "if pressed" lines are ammunition, not part of the first answer. **Golden rule: never claim more than the badge on screen claims.**

### THE AI ITSELF

#### Q · Which LLM are you using? Which model do you prefer?

**The system is deliberately model-agnostic.** Every AI call goes through one module; the provider is a configuration switch — Anthropic, OpenAI, or any OpenAI-compatible internal endpoint, which matters if your locked-down environment supplies its own model. Any frontier model handles intake and planning; what makes the difference is what surrounds the model: every call is grounded in the target repository's own architecture file, rehearsed calls are recorded and replayed deterministically, and prompts are structured so caching keeps cost down.

**If pressed for a name:** don't endorse a vendor in the room — "we've run it against multiple providers; the architecture makes the vendor swappable, and model access inside your environment is a platform approval we've flagged as an open item."

#### Q · Was what we just watched real AI?

**Point at the badge and answer with it.** This environment is simulated end to end and every artifact says so on its face — that labelling discipline is deliberate. Live mode runs the same intake, clarification and planning as real model calls against real GitHub repositories; it has been run end to end, including a real failing CI baseline and real merged pull requests.

**If they want proof:** offer live mode as the encore after the session, not mid-flow. A deterministic run five times out of five beats an impressive one four times out of five.

#### Q · How do you stop hallucination / can we trust the output?

**Three structural answers, not a promise.** First, grounding: every call reads the target repo's own architecture file — components, data model, and explicitly what is *not* part of the application. Second, separation: no phase approves its own output — generation and review are different roles, and a blocked review is never overwritten. Third, gates: a named human signs every decision that matters, and the artifact trail proves who signed what, when, against which version.

#### Q · What does a run cost?

**We log what the provider reports and invent nothing.** Token counts and cache reads are recorded per call; where the provider reports no number, the field stays blank rather than being estimated. Cost per release is a named KPI — we'll quote it from our own measured runs, not from a projection.

### SCOPE & COVERAGE

### Q · How much of the SDLC does the AI actually cover?

**We don't claim 100%, on purpose.** The assessment routes work across streams — frontend, services, data, document intake, mainframe — and labels each task agentic, AI-assisted, or manual. A mainframe field change is manual work that other streams wait on, and the plan says so. An articulated 40–70% survives questioning; a claimed 100% does not.

### Q · Does it write the production code?

**No — by design. S7 is the governed control plane, not an IDE.** Developers build in their own tools — human controlled, AI assisted. S7 generates the governed briefing each team works from, collects evidence from Git and CI rather than status meetings, and runs independent review. The industry's own finding this year is that the bottleneck has moved from writing code to review, quality and governance — that's the part we built.

**If they've seen a Devin-style demo:** "those tools execute inside their own ecosystem; this is the neutral layer that governs an AI agent, a developer in an IDE, and a manual mainframe team the same way — which is what a mixed estate actually needs."

### Q · Where do the estimates come from?

**Today, placeholders — and I'd rather tell you than have you find out.** The intended grounding is historical delivery data: past stories and what they actually took. Estimation accuracy is itself one of the delivery KPIs this is measured on, so the system is designed to be caught out and corrected.

### Q · Won't all these human gates slow delivery down?

**Gates are where a workflow starts, not where it ends.** Approval-on is the starting posture; a workflow earns the removal of a checkpoint by proving itself over many runs. That makes autonomy a measurable maturity ladder instead of a leap of faith — and the run records are exactly the evidence you'd use to climb it.

## FIT & MATURITY

### Q · Can this run in our environment?

**It was built for a locked-down environment from day one.** Plain Python, pinned dependencies, no cloud-managed services, no containers required, and a fresh clone runs the full demonstration offline with zero API keys. The only external dependency any AI feature needs is whichever model endpoint your platform team approves.

### Q · How does this compare to Copilot / GitLab Duo / Rovo / Devin?

**Different layer.** Those are strong execution tools inside their own platforms. None of them is a vendor-neutral control plane spanning multiple teams, multiple repositories and non-AI streams with a client-facing coverage model. The handoff format we publish is the industry-standard one — the same convention GitHub Copilot's own coding agent reads — so those tools plug *into* this rather than compete with it.

### Q · Is this a product we could deploy tomorrow?

**It's a working operating model, demonstrated — not a finished product.** The architecture is the industry's current direction: spec-driven handoff, evidence-based status, explicit human governance. What a funded build adds is depth — tracker integration, real static analysis behind the quality gate, estimation from your historical data. The honest sentence: the skeleton is proven; the flesh is scoped.

### Q · Can we try it with our own requirement document, right now?

**Yes — in live mode, where extraction genuinely reads your document.** In this simulation environment the downstream stages are seeded to the rehearsed scenario, so I won't pretend your document drove them — that's precisely the mislabelling this system exists to prevent. Happy to run a live extraction after the session.

**Questions never to bluff:** production readiness, security certification, real client data handling, and any number (cost, coverage, accuracy) not measured on our own runs. The credibility of every badge depends on never overclaiming once.

## BACK-POCKET MARKET FACTS

- 2026 industry consensus: the SDLC bottleneck has shifted **from code generation to review, quality, governance and architecture** — the layer S7 is.
- Explicit human-in-the-loop governance — who acts alone, who approves, who is accountable — is now the stated enterprise requirement, not a differentiator we invented.
- Spec-driven development is mainstream (GitHub Spec Kit, AWS Kiro, BMAD); versioned specs referenced-not-copied is exactly our delivery-pack model.
- The **AGENTS.md** convention our packs publish is read natively by mainstream coding agents, including GitHub Copilot's.

---

Sources on file: axify.io · tblocks.com · practiq.tech · marktechpost.com · dev.to · github.com/github/spec-kit · ones.com · analyticsinsight.net — retrieved 2026-08-11. MapleSure Insurance is fictional; all demo data is synthetic.