

One requirement. Zero ungoverned steps.

S7 is the delivery scope: a business-driven project carried end to end — from business requirement through design, build, test and production release — where AI does the structuring and a *named human makes every call that matters*.

 FIVE GATES, NONE SKIPPABLE

 NO PHASE SELF-APPROVES

 EVERY ARTIFACT CARRIES ITS PROVENANCE

 RUNS OFFLINE FROM A FRESH CLONE

BUILT BY TCS



S7

What S7 covers, stage by stage

SLIDE 02 · FIVE STAGES · FIVE GATES · ONE COVERAGE MODEL

[Click for the data behind the demo](#)

01



Intake

- Document in, not retyped
- Routed before spend
- Rules captured

02



Planning

- Stories, then tasks
- Criteria carry ids
- Waves and teams

03



Build & Review

- Context published
- Humans implement
- Reviewed independently

04



Quality

- Red baseline first
- Evidence per criterion
- Conditions, not a score

05



Release

- Rendered from records
- Portable document
- Merge stays human



G0

Business Owner
INTAKE GATE



G1

Business Owner
PLAN SIGN-OFF



G2

Independent reviewer
ISOLATED REVIEW



G3

QA Lead
QUALITY GATE



G4

Release Manager
RELEASE APPROVAL

CROSS-CUTTING

The coverage model

The honest answer to "what does the AI actually cover?" — and the seam where the delivery scope hands work to the support scope.



Agentic
enters the automated lane



AI-assisted
externally owned



Manual
hand-work, counted as such

PROVIDED BY THE CLIENT

Requirement document

One

A multi-sprint, business-driven project brief — and nothing else. No application inventory, no ticket extract, no user stories, no data model.

MapleSure-Disability-Claim-Submission-Requirement.pdf

Production records, PII or client-identifiable data, at any point

READ OUT OF THAT DOCUMENT

Nobody retyped any of this

EPIC TITLE	Online Disability Claim Submission for Plan Sponsors
REFERENCE	REQ-2026-114 epic EPIC-S7-001 · target release 2026.R4
OWNER	Group Benefits Operations domain: Group Disability · priority: High
SOURCE	Uploaded PDF, read once requested 2026-07-28 · no transcription step

7
numbered requirements

7
business rules

2
clarifications raised

2
applications affected


LIVE_AI

Measured from run S7-00004. A fresh clone replays the same call offline.

ASSUMPTION

Estimates are effort-weighted placeholders today.

The forward answer is historical delivery data — past stories and the time they actually took. We say that out loud rather than let a number imply otherwise.



Ledger

PROVENANCE

Every artifact

HUMAN · RULE_BASED · LIVE_AI · REPLAYED_AI

GATES HELD

5 of 5

none skippable, all role-separated

OVERRIDES

Recorded

reason mandatory, badged on the workspace


COVERAGE

Articulated

agentic vs. AI-assisted vs. manual

THE SEAM A task small enough goes down the **S3 enhancement lane**; larger chunks are categorised as manual and counted as such — that is where the delivery scope hands off to the support scope.

CONTROL CENTRE



NAMED HUMAN ROLES

- Business Owner — signs G0, G1
- Product Analyst — analysis, plan
- Delivery Lead — orchestration
- Engineering Lead — architecture
- QA Lead — test plan, G3
- Independent Reviewer — G2
- Release Manager — G4
- Support Lead — handover

Role permissions



Phase machine



Provenance labels



STAGE ORCHESTRATION · ONE ENGINE END TO END · ♦ MODEL-BACKED, THE REST DETERMINISTIC

01 Intake workers

- ♦ Requirement Extractor
- ♦ Routing Verdict
- ♦ Intake Analyst + Clarification Raiser
- ♦ Business Rule Extractor
- ♦ New-Application Scaffolder

02 Planning workers

- ♦ Epic-to-Stories Planner
- ♦ Acceptance Criteria Writer
- ♦ Task Splitter — stream, team, repo
- ♦ Dependency Wave Mapper
- ♦ Coverage Defect Corrector

03 Architecture workers

- Architecture Pack Assembler
- ♦ Architecture Refiner — lead proposes
- Architecture Checks
- Version + acceptance tracker

04 Delivery-context workers

- Delivery Pack Renderer
- AC Test Skeleton Generator
- ♦ QA Amendment Refiner
- Git Publisher — AGENTS.md + .s7/**
- Workspace Provisioner

05 Evidence workers

- Git Evidence Sync
- CI Result Joiner — per AC
- ♦ Independent Review Executor
- Staleness Walker
- Quality Handoff Conditions

06 Release workers

- Release Document Renderer
- Approvals Ledger Compiler
- Artifact Exporter — portable packages
- Run Ledger / KPI Reporter

DETERMINISM

Replay by default; the cache key hashes the prompt, so editing a prompt always misses

RESPONSIBLE AI

Staged output is labelled wherever it is shown; a heuristic is never presented as AI

COST CONTROL

Stable cached prefix; deterministic renders instead of calls. Cache counters logged, cost left unset rather than invented

ONE LLM MODULE → PROVIDER OF RECORD



SYSTEMS OF RECORD

< > Git — target repositories

 CI — workflow runs, ci-summary.json

 Artifact plane — deterministic paths

 Run store — ledger + approvals

 Grounding — architecture.md per repo

STRATEGIC DIFFERENTIATORS

Three commitments that survive a hard question

S7 claims coverage of every SDLC deliverable, not depth in any one of them. What convinces a room is how fabrication is stopped and how many stop gates exist — so those are the three things this build spends its effort on.

01

A governed control plane, not another IDE

- Developers keep their own IDE, CLI and Git — S7 never tries to be the editor.
- What S7 owns is the context: generate it, version it, publish it, collect the evidence back.
- Handoff is a file at a deterministic path validated against a schema — not a chat transcript, so nothing is lost at the surface switch.

02

No phase self-approves

- Five gates with separated roles, enforced in the engine — an out-of-order action is refused, not tolerated.
- The generator never accepts its own blueprint; QA never approves tests it did not review; the reviewer cannot edit what it reviews.
- The one escape hatch — unblocking a story early — demands a reason and lands in the approvals ledger.

03

Honest by construction

- Every artifact carries how it was produced — human, rule-based, live AI or replayed AI — and the label ships with it into the document.
- Numbers we have not measured are shown as not measured; cost is left unset rather than estimated.
- Replay-first determinism means the demo that convinced the room is the demo that runs again tomorrow.

THE ESTATE BOUNDARY

What S7 never touches

- Writes only governed context

AGENTS.md · .s7/** · governed test roots

On a fresh s7/... branch

never a default branch

Refuses foreign content

never overwrites what it did not write

Fully reversible

delete the branch – nothing remains

Merging is a human act

never automated, in any mode

WHAT THE ROOM TAKES AWAY

- Requirement to release, one lane
 - Fabrication stopped at the gate
 - A gate you can earn the removal of, once a workflow has proven itself

- Traceability as a field
 - Coverage stated, not claimed

S7 · 01

Intake

S7 · 02

Planning

S7 · 03

Build & Review

S7 · 04

Quality

S7 · 05

Release

S7 · 01

STAGE VALUE & DEPLOYMENT READINESS

Intake

KEY DIFFERENTIATORS

Reads the real document — title, objective and numbered requirements come from what was actually written, never from a canned sample.

Routes before it spends — a cheap verdict decides whether the requirement fits a connected repository or needs a new application.

The analysis raises its own questions — ambiguity becomes a clarification round addressed to the business, not a developer's guess.

Human rules sit beside AI rules — BR-H ids carry human provenance and survive re-analysis; AI extractions stay immutable.

OUTCOME

A structured, grounded epic the business has signed.

PREREQUISITES & TOKEN POSTURE

CONNECTIVITY

Control Centre · LLM module (replay / record / live) · Git remote for repo connection

DATA PREREQUISITES

A requirement document, and zero or more connected repositories, each carrying architecture.md

PROVENANCE EMITTED

RULE_BASED in simulation · LIVE_AI / REPLAYED_AI in live · HUMAN for rules and answers

- TOKEN ESTIMATE – WHY NOT YET

Model access is still pending approval, so no live run has been metered. We log what the provider actually returns and leave cost unset rather than publish a figure we did not measure.

Basis once metered: requirement length + the repo context pack — not the repository

Grounding is a bounded context pack per repo, never a whole codebase in the prompt

The routing verdict runs before analysis, and short-circuits with no model call at all when nothing is connected

The clarification round is capped, not open-ended

S7 · 01
Intake

S7 · 02
Planning

S7 · 03
Build & Review

S7 · 04
Quality

S7 · 05
Release

S7 · 02STAGE VALUE & DEPLOYMENT READINESS

Planning

KEY DIFFERENTIATORS



Traceability is a field, not a paragraph — every acceptance criterion has an id, and tasks point back at the criteria they satisfy.



Coverage is enforced, not hoped for — unclaimed business rules, thin criteria, off-roster teams and dangling dependencies are collected and sent back in one corrective pass.



Waves, teams and repositories — the plan knows who does what, where, and in which order.



Sign-off is load-bearing — it locks the plan and authorises context generation; it does not authorise anyone to write code.

OUTCOME

A signed delivery plan: stories, tasks, criteria and waves.

PREREQUISITES & TOKEN POSTURE

CONNECTIVITY

Control Centre · LLM module · plan cache keyed on rule ids so a new rule forces a fresh call

DATA PREREQUISITES

The signed epic, the merged business-rule set, and the connected repository roster

PROVENANCE EMITTED

LIVE_AI / REPLAYED_AI for the plan · HUMAN for edits, revisions and sign-off

TOKEN ESTIMATE – WHY NOT YET

 Same reason: unmetered until live runs are recorded. Quoting another team's measured figures as ours would be both a confidentiality breach and exactly the mislabelling we refuse elsewhere.

- Basis once metered: epic + merged rule set + repo context, plus at most one corrective call
- One bounded retry names every repairable defect at once, instead of N blind retries
- The plan cache is keyed on rule ids — a re-run with an unchanged rule set costs nothing
- Stage output lands on the artifact plane, so nothing is re-sent to re-plan

S7 · 01

Intake

S7 · 02

Planning

S7 · 03

Build & Review

S7 · 04

Quality

S7 · 05

Release

S7 · 03

STAGE VALUE & DEPLOYMENT READINESS

Build & Review

KEY DIFFERENTIATORS

Human Controlled · AI Assisted — developers implement in their own tooling; S7 generates and publishes the governed context they work from.

Versioned, accepted architecture — immutable v1, v2, v3 directories; a lead proposes a change, the system refines it, acceptance always resets.

Publication is narrow by design — only AGENTS.md, .s7/** and governed test roots, on a fresh branch that can never be a default branch.

Starting work is dependency-gated — per story, never per team, with a recorded override as the only way past.

OUTCOME

Published context, provisioned workspaces, reviewed increments.

PREREQUISITES & TOKEN POSTURE

CONNECTIVITY

Git remote (branch push) · CI provider · Control Centre · LLM module for refinement

DATA PREREQUISITES

The signed plan, the accepted architecture pack, and each target repo's bootstrap record

PROVENANCE EMITTED

RULE_BASED for skeletons and packs · LIVE_AI / REPLAYED_AI for refinement · HUMAN for acceptance, QA approval and assignment · SIMULATED for pseudo-commits

- TOKEN ESTIMATE – WHY NOT YET

Unmetered until live runs are recorded — and most of this stage never calls a model at all, so one headline figure would misrepresent it.
- Basis once metered: architecture pack size plus the number of refinement turns
- Packs, test skeletons and publication are deterministic renders — zero tokens
- Thin task packs reference the plan and architecture by version instead of copying them
- Refinement folds one lead-authored section, not the whole document

S7 · 01

Intake

S7 · 02

Planning

S7 · 03

Build & Review

S7 · 04

Quality

S7 · 05

Release

S7 · 04

STAGE VALUE & DEPLOYMENT READINESS

Quality

KEY DIFFERENTIATORS

A real red baseline — one deliberately-failing test per acceptance criterion, stack-aware, shipped with the pack so CI goes red on purpose first.

Evidence joins back to the criterion — per-test results are matched by name, so a green run means named criteria passed, not that a pipeline finished.

Conditions, never a score — the quality handoff names what is outstanding; a number would hide it.

Staleness rides the provenance walk — revise the architecture and the packs, workspaces and evidence downstream of it are marked stale automatically.

OUTCOME

Per-criterion evidence and a named quality decision.

PREREQUISITES & TOKEN POSTURE

CONNECTIVITY

CI provider (workflow runs + ci-summary.json) · Git remote · Control Centre

DATA PREREQUISITES

Published packs with AC-derived test names, and CI runs from the target repositories


PROVENANCE EMITTED

RULE_BASED for skeletons and joins · HUMAN for the QA gate decision

TOKEN ESTIMATE – NONE TO MAKE

There are no model calls in this stage. Skeletons, CI joins and handoff conditions are deterministic, so there is nothing to meter — and nothing to spend.





- Test skeletons are rule-based and stack-aware — badged RULE_BASED, never sold as AI output
- CI results join to criteria by test name; the match is code, not a model
- Staleness rides the existing provenance walk rather than re-deriving state



STAGE VALUE & DEPLOYMENT READINESS

Release

KEY DIFFERENTIATORS



An auditable release record the business can keep.





Control Centre · Git remote for the delivery branch · no model call required



The completed run: gates, approvals, reviews, evidence and artifacts



RULE_BASED for the document · HUMAN for release approval



- Rendered from stored approvals, verdicts and criteria — zero tokens
- Markdown and HTML come from one renderer, not two separate generations
- Badged RULE_BASED in every mode, including the presenter demo environment

HUMAN DECISIONS

Ten places a human must act

FIVE NUMBERED GATES

- G0

Intake sign-off

Business Owner
- G1

Plan sign-off

Business Owner
- G2

Independent review

Independent Reviewer
- G3

Quality gate

QA Lead
- G4

Release approval

Release Manager

FIVE FURTHER NAMED DECISIONS

- ◆

Architecture acceptance

Engineering Lead · resets on every revision
- ◆

Test-plan approval

QA Lead · an unapproved pack cannot publish
- ◆

Publication

Delivery / Engineering Lead · disposable branch
- ◆

Workspace assignment

Delivery / Engineering Lead · per story
- ◆

Dependency override

Lead only · reason recorded in the ledger

AUTOMATED QUALITY VERIFICATION

Twelve checks, computed — not asked of a model

QC-01 Requirement → story	QC-02 AC → code	QC-03 AC → test
QC-04 Test execution	QC-05 Code coverage	QC-06 Security scan
QC-07 Review gaps	QC-08 Operational readiness	QC-09 Dependency scan
QC-10 Standards check	QC-11 Regression & integration	QC-12 Performance hand-off

⌚ A check that is **arithmetic over recorded evidence cannot hallucinate** — which is why it sits in judgment over the agents, and not the other way round.

THE QUALITY GATE

Three verdicts, in order

- 1

Independent reviewer

finished change vs. acceptance criteria
– not the tests
- 2

Deterministic verdict

nine blocking conditions over the run’s
own records
- 3


Human approval

the final call, with both verdicts in
hand

The nine blocking conditions are QC-01 to QC-08 plus “no required artifact is stale”. The remaining four checks are computed and reported, never silently dropped.

A

UPSTREAM · CONTROL CENTRE — humans decide, agents draft



RECORDED BY EVERY RUN

Evidence the run already holds

✓

Gates held

subject · role · approver · decision · timestamp

📄

Provenance per artifact

author · inputs · version · sha256 · outcome

☰

Criteria with evidence

per-test CI results, joined to the AC by name

⚠️

Overrides

reason mandatory, in the approvals ledger, badged

</>

Publications and amendments

what was written, where, on whose branch

⚡

Per-call telemetry

tokens and cache read/write counters, as the provider reports them

DELIVERY KPIS – S7 IS JUDGED ON THESE, NOT ON SUPPORT KPIS

The seven, and where each one honestly stands

01

Velocity

Stories and tasks completed per sprint

NEEDS DELIVERY HISTORY

02

Cycle time

Stage-to-stage timestamps are already on every ledger entry

RUN-LOCAL ONLY

03

Estimation accuracy

Planned against actual effort per story

ESTIMATES ARE PLACEHOLDERS

04

Defect leakage

Defects found after the quality gate has passed

NEEDS A LIVE RELEASE

05

First-time-right

Stories approved at independent review without a change request

RECORDED PER RUN

06

On-time / on-budget

The signed plan as a baseline, against what actually happened

NEEDS A REAL BASELINE

07

Cost per release

Token spend and effort attributed to one release

COUNTERS LOGGED, COST UNSET

🕒

Every row above is either recorded by a run today or marked as not yet measured. Nothing on this slide is estimated to fill a gap — that is the same rule the artifacts follow.

CONSOLIDATED SCORECARD

Efficiency

S1-S6

Effort reduction · productivity per FTE

S7

Cycle time · cost per release

Service quality

S1-S6

SLA / XLA adherence

S7

First-time-right · gates held without override

Issue resolution

S1-S6

MTTR · reopen and escalation rates

S7

Defect leakage after release

Delivery productivity

S1-S6

Backlog ageing

S7

Velocity · estimation accuracy · on-time delivery



S7

What a live run needs here

SLIDE 11 · FOUR DEPENDENCIES · NONE OF THEM CHANGE THE PIPELINE

Questions

CLOSING

Every gate already holds. What it needs next is access.

BUILT BY TCS

Simulation and replay run the whole lane offline from a fresh clone today — five gates, provenance on every artifact, and a release document rendered from the run's own records. The four items below are what turn that same run into a live one in your environment. **None of them changes the shape of the pipeline.**



Model access

An approved internal assistant, a self-hosted gateway, or a key. Every call goes through one module, chosen by an environment variable.

Simulation and replay need none of it



A port, or the CLI

The Control Centre is static files bound to 127.0.0.1 — no CDN, no runtime install, no container. It is the surface where a human decides.

If a port cannot be served, the CLI carries the same ledger



A domain owner

The submission scenario needs an SME to confirm forms, required attachments, status names and pre-population rules.

Scenario detail, not architecture



Delivery history

Past stories and the time they actually took — the named grounding source for estimation, and the only honest one.

Until it lands, estimates stay labelled placeholders